

Part 1

On Shopify, we have exactly 100 sneaker shops, and each of these shops sells only one model of shoe. We want to do some analysis of the average order value (AOV). When we look at orders data over a 30 day window, we naively calculate an AOV of \$3145.13. Given that we know these shops are selling sneakers, a relatively affordable item, something seems wrong with our analysis.

We are to determine using three questions if the calculation of AOV is accurate and if not which better metric could we use.

Importing the necessary libraries

```
import matplotlib as mpl
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
sns.set_style("darkgrid")
mpl.rcParams['figure.figsize'] = (20,5)
```

Reading the CSV file

This data set is publicly provided in the shopify data science intern job posting and has been downloaded and saved as a csv file

Reading the csv file

```
shop_stores = pd.read_csv(r'C:\Users\Public\Documents\Shopify
Application\Data Science Intern Challenge Data Set - Sheet1.csv')
```

After our data containing the shopify stores have been read, The next step is to visually inspect the data and make sure it has been read properly. The first 10 rows of the data would be printed for the visual inspection. We would also check for null values and make sure the data types for the columns of interest are read properly

print out the first 20 rows

```
shop_stores.head(10)
```

	order_id	shop_id	user_id	order_amount	total_items
0	1	53	746	224	2
1	2	92	925	90	1
2	3	44	861	144	1
3	4	18	935	156	1
4	5	18	883	156	1

5	6	58	882	138	1
credit_card					
6	7	87	915	149	1
cash					
7	8	22	761	292	2
cash					
8	9	64	914	266	2
debit					
9	10	52	788	146	1
credit_card					

	created_at
0	2017-03-13 12:36:56
1	2017-03-03 17:38:52
2	2017-03-14 4:23:56
3	2017-03-26 12:43:37
4	2017-03-01 4:35:11
5	2017-03-14 15:25:01
6	2017-03-01 21:37:57
7	2017-03-08 2:05:38
8	2017-03-17 20:56:50
9	2017-03-30 21:08:26

Printing the information from the imported file to perform quick sanity check

```
shop_stores.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 5000 entries, 0 to 4999
```

```
Data columns (total 7 columns):
```

#	Column	Non-Null Count	Dtype
0	order_id	5000 non-null	int64
1	shop_id	5000 non-null	int64
2	user_id	5000 non-null	int64
3	order_amount	5000 non-null	int64
4	total_items	5000 non-null	int64
5	payment_method	5000 non-null	object
6	created_at	5000 non-null	object

```
dtypes: int64(5), object(2)
```

```
memory usage: 273.6+ KB
```

checking how many unique entries there are in the data

```
print(len(pd.unique(shop_stores['shop_id'])))
```

```
100
```

From our sanity check

1. There are 5000 data points
2. There are 100 unique stores that sell sneakers

3. There are no null values
4. The data type for all columns are properly attributed except the payment method and created at columns that have the data type object instead of string and datetime64. For the purpose of this analysis, these columns would not be converted as they are not needed

Question 1. Think about what could be going wrong with our calculation of AOV? Think about a better way to evaluate this data.

Answer Summary

The average order value of 3145.13 is skewed because of two assumptions that have been tested

1. Some stores carry sneakers that are significantly more expensive than orders. Infact at least one store sells one sneaker for 25,725 dollars
2. Two stores generally have higher orders than the rest and these two stores are shop id 42 and 78.

With that being said, the average or mean would not be considered the better measure of central tendency

Below is how I arrived at the answer

```
# Describing the order amount column
shop_stores.order_amount.describe()
```

```
count      5000.000000
mean       3145.128000
std        41282.539349
min         90.000000
25%        163.000000
50%        284.000000
75%        390.000000
max       704000.000000
Name: order_amount, dtype: float64
```

As shown above, the average order is calculated as 3145.13. However, two other calculations stand out and these are the minimum order and the maximum order. The minimum order amount is 90.00 and the maximum as shown is 704,000. This is such a large difference. One assumption could be that some stores carry sneakers that are way more expensive than others. To check this assumption, we would create a caluclated column to show the price per item sold.

Creating a calculated column for price per unit for the stores.

```
# Creating the column price per unit
```

```
shop_stores['price_per_unit'] = shop_stores['order_amount'] /
shop_stores['total_items']
```

```
shop_stores.head()
```

	order_id	shop_id	user_id	order_amount	total_items
0	1	53	746	224	2
1	2	92	925	90	1
2	3	44	861	144	1
3	4	18	935	156	1
4	5	18	883	156	1

	created_at	price_per_unit
0	2017-03-13 12:36:56	112.0
1	2017-03-03 17:38:52	90.0
2	2017-03-14 4:23:56	144.0
3	2017-03-26 12:43:37	156.0
4	2017-03-01 4:35:11	156.0

Now that the Price per unit column has been created, the next step would be to describe the column now to check for the average prices in the store

```
# what is the average prices in the store?  
shop_stores.price_per_unit.describe()
```

```
count      5000.000000  
mean       387.742800  
std        2441.963725  
min         90.000000  
25%        133.000000  
50%        153.000000  
75%        169.000000  
max       25725.000000  
Name: price_per_unit, dtype: float64
```


As we can see, on average the sneakers sell for 387 dollars. However, again pointing out that there is at least one store that sells the sneakers for 90 dollars and at least one store that sells sneakers for 25,725 dollars. So the assumption that some stores have sneakers that are way more expensive is considered true. This could skew our calculation of Average order Value.

The next assumption to be considered is that the orders are significantly higher for some stores than others. To check for this assumption (or hypothesis) we would plot a scatter plot to show the order amounts versus the shop id.

```
# Plot a scatterplot using for total items vs the order amount
```

```
sns.scatterplot(data = shop_stores, x = 'shop_id', y = 'order_amount')
```

```
<AxesSubplot:xlabel='shop_id', ylabel='order_amount'>
```



As shown in above, there are some stores that have significantly higher orders than others. This could also potentially skew the results for the average order value. The next step would be to find which stores have order values higher than the average amount of 3145.13

```
store_over_average= shop_stores[shop_stores['order_amount'] > 3145.13]  
print(store_over_average['shop_id'])
```

```
15      42  
60      42  
160     78  
490     78  
493     78
```

```
..  
4646    42  
4715    78  
4868    42  
4882    42  
4918    78
```

```
Name: shop_id, Length: 63, dtype: int64
```

Our second assumption tests true. There are two stores that significantly have higher orders than the rest and those stores have the shop id 42 and 78. Again, these higher order amounts in these two stores would definitely skew our average order value

Question 2. What metric would you report for this dataset?

Answer Summary to Question 2.

Removing the outliers would not be a very good idea based on the explanation given above. There are three measures of central tendencies. Mean, median and mode. The mean generally would be a good measure however when there are data points that skew our data, the median would be considered the better measure of central tendencies. The median shows the mid point of the dataset and is not influenced by the outliers.

Due to our findings, we have established that the mean is not the better measure of central tendency in our data.

In order to determine a more accurate average order value we could remove the outliers and measure the average again. Let's take a look at the Outliers.

Should we remove the Outliers and recalculate the AOV?

To identify the Outlier, we need to

1. Define the quartiles Q1 and Q3.
2. Calculate the Interquartile Range ($Q3 - Q1$)
3. Create two new variables; Lower_Range ($Q1 - 1.5 \times IQR$) and Upper_Range ($Q3 + 1.5 \times IQR$)
4. Filter the DataFrame for outliers and remove them

After, filtering the dataframe for outliers and removing them then we would Calculate what the proportion of outliers exist (i.e. Number of entries left after outlier removal / Number of total entries in dataset)

Generally an outlier is defined as any Value that is either $Q1 - 1.5 \times IQR$ or greater than $Q3 + 1.5 \times IQR$.

(i) Creating the 25th and 75th percentile

```
Q1 = shop_stores.quantile(0.25)
```

```
Q3 = shop_stores.quantile(0.75)
```

(ii) Calculate the interquartile range

```
IQR = Q3 - Q1
```

```
print(IQR)
```

```
order_id      2499.5
shop_id        51.0
user_id       150.0
order_amount   227.0
total_items     2.0
price_per_unit  36.0
dtype: float64
```

(i) Define two new variables, Lower_Limit and Upper_Limit

```
lower_limit = Q1 - 1.5 * IQR
```

```
upper_limit = Q3 + 1.5 * IQR
```

Filter the raw dataframe to include only the outliers

```
outliers = shop_stores[((shop_stores < lower_limit) | ((shop_stores > upper_limit)))].any(axis=1)]
```

Percentage of data left after we remove outliers

```
outliers_percent = 100 * len(outliers) / len(shop_stores)
```

```
shop_stores_percent_left = 100-outliers_percent
```

```
print(shop_stores_percent_left)
```

96.62

```
<ipython-input-12-74216d9d2a86>:7: FutureWarning: Automatic reindexing
on DataFrame vs Series comparisons is deprecated and will raise
ValueError in a future version. Do `left, right = left.align(right,
axis=1, copy=False)` before e.g. `left == right`
```

```
outliers = shop_stores[((shop_stores < lower_limit) | ((shop_stores
> upper_limit))).any(axis=1)]
```

```
<ipython-input-12-74216d9d2a86>:7: FutureWarning: Automatic reindexing
on DataFrame vs Series comparisons is deprecated and will raise
ValueError in a future version. Do `left, right = left.align(right,
axis=1, copy=False)` before e.g. `left == right`
```

```
outliers = shop_stores[((shop_stores < lower_limit) | ((shop_stores
> upper_limit))).any(axis=1)]
```

We have 96.62% of our data left. While this is still considered a good amount of data points left, we miss out on important information. For example our calculation would exclude specifically two shops with ids 42 and 78

Question 3 and Answer to What is its value?

Calculate the median of the order amount

```
print("The better measure of central tendency is the median and the
median order value is")
```

```
print(shop_stores['order_amount'].median())
```

The better measure of central tendency is the median and the median
order value is
284.0

Part 2

a) How many orders were shipped by Speedy Express in total?

```
SELECT ShipperName, COUNT(0.ShipperID) AS Orders_Shipped
FROM Orders o
INNER JOIN Shippers s
ON o.ShipperID = s.ShipperID
WHERE ShipperName = 'Speedy Express'
Group by o.ShipperID
```

Speedy Express had 54 orders shipped in total

b) What is the last name of the employee with the most orders?

```
SELECT o.EmployeeID, e.LastName, COUNT(o.OrderID) AS Total_Orders
FROM Orders o
INNER JOIN Employees e
ON o.EmployeeID = e.EmployeeID
GROUP BY o.EmployeeID
```

```
ORDER BY Total_Orders DESC  
LIMIT 1
```

The last name of the employee with the most orders is Peacock and total orders from Peacock is 40

c) What product was ordered the most by customers in Germany?

```
SELECT sum(od.Quantity) as Total_Orders, p.ProductName FROM  
OrderDetails od  
LEFT JOIN Products p ON p.ProductID = od.ProductID  
LEFT JOIN Orders o ON o.OrderID = od.OrderID  
LEFT JOIN Customers c ON c.CustomerID = o.CustomerID  
WHERE c.Country = 'Germany' GROUP BY p.ProductName ORDER BY  
Total_Orders DESC  
LIMIT 1
```

The Boston Crab meat was ordered the most by customers in Germany